

SPoP Project

Dr. Rahul Jain, rENIAC Inc

Problem

- Build an In-Memory Key-Value Storage Software in C++
- Supported APIs
 - `get(key)`: returns value for the key
 - `put(key, value)`: add key-value, overwrite existing value
 - `delete(key)`
 - `get(int N)`: returns Nth key-value pair
 - `delete(int N)`: delete Nth key-value pair
- Spec
 - Max key size: 64 bytes
 - Each key char can be (a-z) or (A-Z): Total 52 possible chars
 - Max Value Size: 256 bytes, any ASCII value
 - No DS/boost/STL etc Libraries to be used

Other Spec

- Only pthread to be used for multithreading
- Keys to be in ~~Lexicographical~~ strcmp (cstring) Order
- API calls would be blocking



```
#include <iostream>
#include <cstring>
using namespace std;
int main(int argc, char** argv) {
    if(argc !=3) cout << "Usage lex str1 str2" << endl;
    else cout << "strcmp(" << argv[1] << ", " << argv[2] << ") = " << strcmp(argv[1], argv[2]) << endl;
}
```

Class Definition

```
struct Slice{
    uint8_t size;
    char*    data;
};

class kvStore {
public:
    kvStore(uint64_t max_entries);
    bool get(Slice &key, Slice &value): //returns false if key didn't exist
    bool put(Slice &key, Slice &value): //returns true if value overwritten
    bool del(Slice &key);
    bool get(int N, Slice &key, Slice &value): //returns Nth key-value pair
    bool del(int N): //delete Nth key-value pair
};
```

Evaluation Benchmark

- Multithreaded benchmark application
- Runs:
 - Benchmark would first load data via put calls (10 million entries, 2 min time limit)
 - Perform Single Threaded Transactions to verify kvStore functionality
 - **Multiple Transaction Threads with each thread calling one of the APIs**
- Evaluation Metrics:
 - TPS (Transactions Per Second)
 - **Average** CPU Usage
 - **Peak** Memory Usage

Ranking Methodology (Might Change)

- Normalize all 3 metrics with Avg
- $\text{Score} = (\text{TPS} * \text{TPS}) / (\text{CPU} * \text{Mem})$

Relevant Concepts

- Ordered Data Structures
 - B-Tree, BST, etc
 - Tries, FST (Advanced)
 - Hybrid: Combination of Hash Table + Tree
 - <https://www.coursera.org/learn/cs-fundamentals-2>
- Bit Hacks
- Cache Optimizations
- Memory Allocation Optimizations
 - Memory Object Reuse
 - No Dynamic Memory Allocation
- Multithreading (pthreads)
- These are only starting pointers, feel free to explore

Submission Deadline

- Form Groups of 2, can be different from assignment group
 - No Groups of 3 allowed. Group of 1 is OK
- Read papers or other relevant material around better data structures
 - **must provide reference**
- Implementation Spec: 3rd Feb
- Code+Report: 18th Feb (**Hard Deadline**)
- If a working code has been submitted on/before 18th Feb, a more optimized version can be submitted upto **29th Feb**